



ZJU-UIUC Institute

Zhejiang University / University of Illinois at Urbana-Champaign Institute



ECE 470: Introduction to Robotics

Lecture 25

Liangjing Yang

Assistant Professor, ZJU-UIUC Institute

liangjingyang@intl.zju.edu.cn

Wechat ID: Liangjing_Yang

Overview of Robot Vision

O. Introduction to Robot Vision

- What is Robot Vision?

I. Image Formation

- The science behind machine vision (+ represent as a form of signal)

II. Image Processing

- Common techniques to manipulate, enhance & analyse images

III. Robot Vision Applications

- 3D Vision; Photogrammetry; Vision-based techniques in robotics- visual servo, pose estimation, localization, mapping, navigation

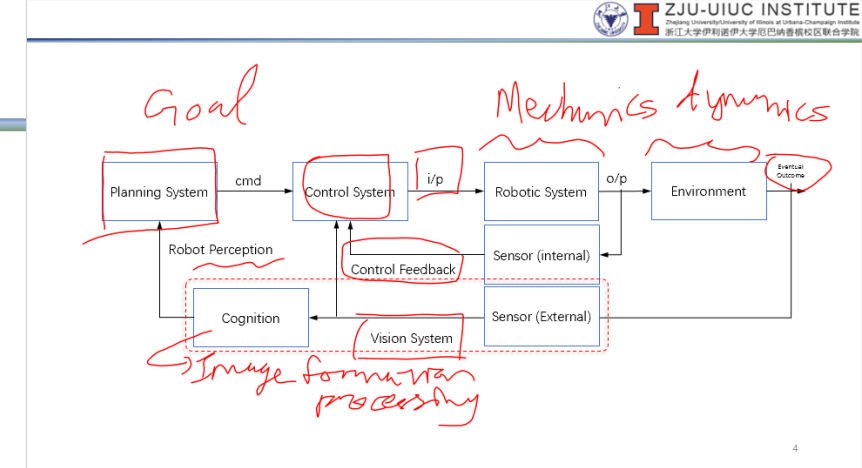





Image Processing

- ☒ Thresholding & Histogram Processing
- ☒ Filtering  
- ☒ Edge & Corner Detection
- Interest Points/ Feature Description 
- Lines & Shapes

Represent
extra (neighborhood info)

Recap: Canny Edge Detection

- Involves 5 steps:
 1. Apply Gaussian filter: remove the noise
 2. Find the intensity gradients of the image
 3. Apply non-maximum suppression: remove spurious response
 4. Apply edge detection using two threshold value
 5. Finalize edge detection by hysteresis



Canny Edge Detection

- Step (1): Apply Gaussian filter to smooth the image in order to remove the noise
 - Edge detection are easily affected by image noise
 - A Gaussian filter is applied to convolve with the image
 - $h(x, y) = \frac{1}{2\pi\sigma^2} e^{-(x^2+y^2)/2\sigma^2}$
 - This will smooth the image to reduce the effects of obvious noise on the edge detector

Canny Edge Detection

- Step (2): Find the intensity gradients of the image


$$\begin{aligned} \frac{\partial I(x, y)}{\partial x} &= \frac{I(x+1, y) - I(x-1, y)}{2} \\ \frac{\partial I(x, y)}{\partial y} &= \frac{I(x, y+1) - I(x, y-1)}{2} \end{aligned}$$

$$\text{Gradient Vector: } \nabla I(x, y) = \left[\frac{\partial I(x, y)}{\partial x}, \frac{\partial I(x, y)}{\partial y} \right]^T$$

$$\begin{aligned} \text{mag} \rightarrow |\nabla I(x, y)| &= \sqrt{\left(\frac{\partial I(x, y)}{\partial x} \right)^2 + \left(\frac{\partial I(x, y)}{\partial y} \right)^2} \\ \text{ang} \rightarrow \theta(x, y) &= \tan^{-1} \left(\frac{\partial I(x, y)}{\partial y} / \frac{\partial I(x, y)}{\partial x} \right) \end{aligned}$$

Canny Edge Detection

- Step (3): Apply non-maximum suppression to get rid of spurious response to edge detection
 - compare values along the gradient vector to retain only the maximum



Canny Edge Detection

- Step (4): Apply edge detection using two threshold value K_H and K_L
 - $\text{Edge}(x, y) = \begin{cases} E_{\text{strong}} & \text{if } |\nabla I(x, y)| > K_H \\ E_{\text{average}} & \text{if } K_L \leq |\nabla I(x, y)| \leq K_H \\ E_{\text{weak}} & \text{if } |\nabla I(x, y)| < K_L \end{cases}$

- Step (5): Finalize edge detection by hysteresis thresholding
 - Small K means more details
 - High K includes more noise
 - Hysteresis Thresholding allows us to apply both
 - Keep both high threshold K_H and low threshold K_L .
 - Any edges with magnitude $< K_L$ are discarded
 - Any edges with magnitude $> K_H$ are kept
 - An edge with magnitude between the two threshold values is kept if there is a path of edges with magnitude $> K_L$ connecting the edge to another edge with magnitude $> K_H$

Recap: Harris Corner Detection

- Involves 5 steps:
 1. Apply Gaussian filter: remove the noise
 2. Find the intensity gradients of the image
 3. Apply non-maximum suppression: remove spurious response
 4. Apply edge detection using two threshold value
 5. Finalize edge detection by hysteresis

Looking at the Detector as an operation

Corner Detection

- Basic idea of corner detection: large change in appearance
 - Flat region: no change
 - Edge: no change along edge
 - Corner: significant change in more than one direction

Images taken from Introduction to Autonomous Mobile Robot

Harris Corner Detector

SSD - sum squared difference

- Consider taking an image patched centered on (u, v) and shifting it by (x, y) , the sum of square differences SSD between these two patches is:

$$SSD(x, y) = \sum_u \sum_v [I(u, v) - I(u + x, v + y)]^2$$
- Using first-order Taylor expansion,

$$I(u + x, v + y) \approx I(u, v) + I_x(u, v)x + I_y(u, v)y$$
- Hence, SSD becomes

$$SSD(x, y) \approx \sum_u \sum_v [I_x(u, v)x + I_y(u, v)y]^2$$

$$= \sum_u \sum_v [I_x^2 x^2 + 2xy I_x I_y + I_y^2 y^2]$$

Handwritten notes: grid, x, y, Ix, Iy, M

Harris Corner Detector

$$SSD(x, y) = \sum_u \sum_v [I_x^2 x^2 + 2xy I_x I_y + I_y^2 y^2]$$

$$= \sum_u \sum_v [x \ y] \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

$$= [x \ y] \sum_u \sum_v \begin{bmatrix} I_x^2 & I_x I_y \\ I_x I_y & I_y^2 \end{bmatrix} \begin{bmatrix} x \\ y \end{bmatrix}$$

- $SSD(x, y) = [x \ y] M \begin{bmatrix} x \\ y \end{bmatrix}$
- Since M is symmetric, we can rewrite the matrix as:

$$M = A^{-1} \begin{bmatrix} \lambda_1 & 0 \\ 0 & \lambda_2 \end{bmatrix} A$$
- where λ_1 and λ_2 are the eigenvalues of M

Harris Corner Detector

non corner

- As mentioned, a corner is characterized by a large variation of SSD in all direction, the larger the variation in that direction
- Both λ are small means flat region
- One strong and one weak λ means edge
- Two strong λ means corner


Quick way of Calculating Corner Response:

$$R = \lambda_1 \lambda_2 - k \cdot (\lambda_1 + \lambda_2)^2 = \det(M) - k \cdot \text{tr}(M)^2$$

where k is an empirically determined constant: $k \in [0.04, 0.06]$.

Feature Extraction

- (Local) Features
 - (Locally) Unique structures e.g. Edges, Corners, Blobs
- Detector
 - Methods for deciding if a pixel is associated with particular a feature e.g. Canny, Harris, LoG
- ✓ • Descriptor
 - Representations of (local) features



Feature Extraction

- Ideal features should be
 - Repeatable
 - Robustly detectable in different viewpoints and imaging conditions (e.g. noise)
 - Distinctive
 - Uniquely representable for comparison
 - Localizable
 - Provide spatial information

Feature Extraction

- Detector
 - Methods for deciding if a pixel is associated with particular a feature e.g. Canny, Harris, LoG

| Detector | Feature Type | Function | Scale Independent |
|----------------------------------|-------------------------------|--|-------------------|
| FAST [1] | Corner | detectFASTFeatures | No |
| Minimum eigenvalue algorithm [4] | Corner | detectMinEigenFeatures | No |
| Corner detector [3] | Corner | detectHarrisFeatures | No |
| SURF [11] | Blob | detectSURFFeatures | Yes |
| KAZE [12] | Blob | detectKAZEFeatures | Yes |
| BRISK [6] | Corner | detectBRISKFeatures | Yes |
| MSER [8] | Region with uniform intensity | detectMSERFeatures | Yes |
| ORB [13] | Corner | detectORBFeatures | No |



Feature Extraction

• Detector

- Methods for deciding if a pixel is associated with particular a feature e.g. Canny, Harris, LoG


| Detector | Feature Type | Function | Scale Independent |
|----------------------------------|-------------------------------|-------------------------------------|-------------------|
| FAST [1] | Corner | <code>detectFASTFeatures</code> | No |
| Minimum eigenvalue algorithm [4] | Corner | <code>detectMinEigenFeatures</code> | No |
| Corner detector [3] | Corner | <code>detectHarrisFeatures</code> | No |
| SURF [11] | Blob | <code>detectSURFFeatures</code> | Yes |
| KAZE [12] | Blob | <code>detectKAZEFeatures</code> | Yes |
| BRISK [6] | Corner | <code>detectBRISKFeatures</code> | Yes |
| MSER [8] | Region with uniform intensity | <code>detectMSERFeatures</code> | Yes |
| ORB [13] | Corner | <code>detectORBFeatures</code> | No |

speed up Robust Feature

Feature Extraction

feature

Block



• Descriptor

- Representations of (local) features


| Descriptor | Binary | Function and Method | Invariance | | Typical Use | |
|--|--------|--|------------|----------|-------------------------------|----------------|
| | | | Scale | Rotation | Finding Point Correspondences | Classification |
| HOG | No | <code>extractHOGFeatures(I, ...)</code> | No | No | No | Yes |
| LBP | No | <code>extractLBPFeatures(I, ...)</code> | No | Yes | No | Yes |
| SURF | No | <code>extractFeatures(I, points, 'Method', 'SURF')</code> | Yes | Yes | Yes | Yes |
| KAZE | No | <code>extractFeatures(I, points, 'Method', 'KAZE')</code> | Yes | Yes | Yes | Yes |
| FREAK | Yes | <code>extractFeatures(I, points, 'Method', 'FREAK')</code> | Yes | Yes | Yes | No |
| BRISK | Yes | <code>extractFeatures(I, points, 'Method', 'BRISK')</code> | Yes | Yes | Yes | No |
| ORB | Yes | <code>extractFeatures(I, points, 'Method', 'ORB')</code> | No | Yes | Yes | No |
| <ul style="list-style-type: none"> • Block • Simple pixel neighborhood around a keypoint | No | <code>extractFeatures(I, points, 'Method', 'Block')</code> | No | No | Yes | Yes |

Feature Extraction Practical Consideration

- Select appropriate types of feature-detector-descriptor
 - Based on
 - • Nature of scenes
 - • Application context
 - • Operation requirement

Feature Extraction Practical Consideration

- Select appropriate types of feature-detector-descriptor
 - Based on
 - • Nature of scenes
 - • Application context
 - • Operation requirement



Feature Matching

- Given views of the same scene, how to match the features?
 - Detect; Describe; Match; Transform (for image mapping) or Tracking (Visual tracking)


Feature Detection, Description,
& Matching

Feature Matches


Detect




Describe



Match





Tracking



Similarity Scores between image patches

Measure

- How do we measure the similarity/ difference?



Similarity Score

- Sum of Squared Difference (SSD)

SSD between two image patches I_1 and I_2

$$w_{SSD}(x, y) = \sum_p^P \sum_q^Q [I_1(p, q) - I_2(p, q)]^2$$

Recall in corner detection, we look at change of patch I centered on (p, q) and itself when shifted by (x, y) , we used SSD to represent the change

$$w_{SSD}(x, y) = \sum_p^P \sum_q^Q [I(p, q) - I(p + x, q + y)]^2$$

Similarity Score

- Sum of Squared Difference (SSD)

SSD between two image patches I_1 and I_2

$$w_{SSD}(x, y) = \sum_p^P \sum_q^Q [I_1(p, q) - I_2(p, q)]^2$$

Recall in corner detection, we look at change of patch I centered on (p, q) and itself when shifted by (x, y) , we used SSD to represent the change

$$w_{SSD}(x, y) = \sum_p^P \sum_q^Q [I(p, q) - I(p + x, q + y)]^2$$

Similarity Score

- Normalized Cross-correlation (NCC)

For a $U \times V$ image and $P \times Q$ patch, the cross-correlation $w_{cc}(u, v)$ at a particular image coordinates (u, v) of a template patch $g(p, q)$ and the image $f(p, q)$ is

$$w_{cc}(u, v) = \sum_{p=0}^P \sum_{q=0}^Q g(p, q) f(p+u, q+v)$$

To reduce the over sensitivity toward intensity variant,
normalized cross-correlation coefficient (Similarity Score),


$$w_{ncc}(u, v) = \frac{\sum_{p=0}^P \sum_{q=0}^Q (g(p, q) - \bar{g})(f(p+u, q+v) - \bar{f}(u, v))}{\left[\left(\sum_{p=0}^P \sum_{q=0}^Q (g(p, q) - \bar{g})^2 \right) \left(\sum_{p=0}^P \sum_{q=0}^Q (f(p+u, q+v) - \bar{f}(u, v))^2 \right) \right]^{0.5}}$$

is used where \bar{g} and \bar{f} denotes the mean intensity in the template and overlapping region, respectively.

Detection of Line and Shape

- After detecting the edges and local interest points, how do we detect lines and other geometries?
- A problem of pattern recognition

circle
triangle
polygons




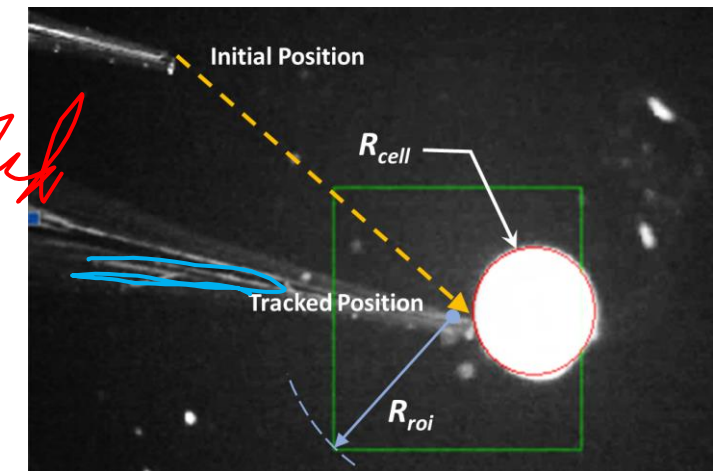

Microscope
Image



Surgical Scene

Detection of Line and Shape

- After detecting the edges and local interest points, how do we detect lines and other geometries?
- A problem of pattern recognition




Circle detection

Fitting lines to edges using least square

~~E~~ C ~~E~~ 470


- Given:
 - Many edges (x_i, y_i)
 - Equation of line: $y = mx + c$
 - Parameters: m, c
- Objective function is to minimize average square distance:
 - $E = \sum \frac{(y_i - mx_i - c)^2}{N}$
- Using $\frac{\partial E}{\partial m} = 0$ & $\frac{\partial E}{\partial c} = 0$



$$c = \bar{y} - m \bar{x}$$

$$m = \frac{\sum_i (x_i - \bar{x})(y_i - \bar{y})}{\sum_i (x_i - \bar{x})^2}$$

Fitting lines to edges using least square




$$y = mx + c$$
$$\begin{bmatrix} y \\ 1 \end{bmatrix} = \begin{bmatrix} m & c \end{bmatrix} \begin{bmatrix} x \\ 1 \end{bmatrix}$$

- Possible Solution: Hough Transformation

Hough Transform


- Ability to detect edges
- What about lines? Shapes (eg: circle)?
- Hough Transform
 - Elegant method for direct object recognition
 - Edges need not be connected
 - Key idea: edges “vote” for the possible model

$\theta = \arctan(\frac{y}{x})$



Hough Space – (ρ, θ) space

- $\rho = x \cos \theta + y \sin \theta$
- $y = \frac{\rho}{\sin \theta} - \frac{x}{\tan \theta}$
- Line equation: $\rho = x \cos \theta + y \sin \theta$
 - Parameters: ρ and θ
 - Where $0 \leq \theta \leq 2\pi$



Hough Space – ρ , θ space


- $\rho = x \cos \theta + y \sin \theta$

- $y = \frac{\rho}{\sin \theta} - \frac{x}{\tan \theta}$ $y = mx + c$

- Line equation: $\rho = x \cos \theta + y \sin \theta$

- Parameters: ρ and θ


- Where $0 \leq \theta \leq 2\pi$



Hough Space



Hough Space – (ρ, θ) space

- Line equation: $\rho = x \cos \theta + y \sin \theta$
 - Parameters: ρ and θ
 - Where $0 \leq \theta \leq 2\pi$
- How to map onto the Hough Space (sometimes also known as parameter space)?
 - $\rho = x_i \cos \theta + y_i \sin \theta$
 - Sinusoid curve



Hough Space – (ρ, θ) space


- $\rho = x \cos \theta + y \sin \theta$
- Points in picture \rightarrow sinusoids in parameter space
- Points in parameter space \rightarrow lines in picture
- There will be a unique intersection point if the points in the picture form a straight line



But how to choose the solution if there are noise?

Hough Space – (ρ, θ) space

- $\rho = x \cos \theta + y \sin \theta$
- Points in picture \rightarrow sinusoids in parameter space
- Points in parameter space \rightarrow lines in picture
- There will be a unique intersection point if the points in the picture form a straight line




But how to choose the solution if there are noise?

Quantizing parameter space and Voting



- Discretize the (ρ, θ) space
 - For each point (x_i, y_i) , compute only for a finite set of angles $\theta = \theta_1, \theta_2, \dots, \theta_N$
 - For each θ_j , obtain $\rho_{ij} = x_i \cos \theta_j + y_i \sin \theta_j$
- Create a matrix, called the accumulator matrix
 - Each column corresponds to angles $\theta = \theta_1, \theta_2, \dots, \theta_N$
 - Each row corresponds to the “bins” (intervals) of the resulting distance ρ
- Voting:
 - For each point in the image and for each θ_j , compute the ρ_{ij} , and increment the corresponding element of the accumulator matrix
 - Highest value means highest “vote”


Quantizing parameter space and Voting

- Voting (continue)
 - If more than one line, you can set a threshold value (number of vote) to obtain more lines
 - For example, if number of votes (ie value in that element) is more than the threshold value, then consider that (ρ, θ) to be a line



Quantizing parameter space and Voting

- Voting (continue)
 - If more than one line, you can set a threshold value (number of vote) to obtain more lines
 - For example, if number of votes (ie value in that element) is more than the threshold value, then consider that (ρ, θ) to be a line  




Case Problem (1)

Using what you have learnt, design an algorithm that is capable of identifying the following simple blocks on white background.

1)Rectangle



2)Triangle



Solution

Case Problem (1) Solution

- Stage (1): **Edge detection** using Canny edge detection
 - Apply Gaussian filter to smoothen the image in order to remove the noise
 - Find the intensity gradients vector of the image
 - $|\nabla I(x, y)| = \sqrt{\left(\frac{\partial I(x, y)}{\partial x}\right)^2 + \left(\frac{\partial I(x, y)}{\partial y}\right)^2}$
 - $\theta(x, y) = \tan^{-1} \left(\frac{\partial I(x, y)}{\partial y} / \frac{\partial I(x, y)}{\partial x} \right)$
 - Apply non-maximum suppression to get rid of spurious response to edge detection
 - Apply edge detection using two threshold value K_H and K_L
 - $Edge(x, y) = \begin{cases} E_{strong} & \text{if } |\nabla I(x, y)| > K_H \\ E_{average} & \text{if } K_L \leq |\nabla I(x, y)| \leq K_H \\ E_{weak} & \text{if } |\nabla I(x, y)| < K_L \end{cases}$ Finalize edge detection by hysteresis
 - Any edges with magnitude $< K_L$ are discarded
 - Any edges with magnitude $> K_H$ are kept
 - An edge with magnitude between the two threshold values is kept if there is a path of edges with magnitude $> K_L$ connecting the edge to another edge with magnitude $> K_H$

Case Problem (1) Solution

- Stage (2): Hough Transformation – **line detection**
 - Use $\rho = x \cos \theta + y \sin \theta$
 - Map it onto the Hough space – (ρ, θ) space
 - Intersection point exist if the points in the picture form a straight line
 - As more than one line is involved, various (ρ, θ) are considered to be a line number if number of votes is more than the threshold value

Case Problem (1) Solution

- Stage (2): Hough Transformation - **voting**
 - Discretize the (ρ, θ) space
 - For each point (x_i, y_i) , compute only for a finite set of angles $\theta = \theta_1, \theta_2, \dots, \theta_N$
 - For each θ_j , obtain $\rho_{ij} = x_i \cos \theta_j + y_i \sin \theta_j$
 - Create a matrix, called the accumulator matrix
 - Each column corresponds to angles $\theta = \theta_1, \theta_2, \dots, \theta_N$
 - Each row corresponds to the “bins” (intervals) of the resulting distance ρ
 - Voting:
 - For each point in the image and for each θ_j , compute the ρ_{ij} , and increment the corresponding element of the accumulator matrix
 - Highest value means highest “vote”

Case Problem (1) Solution

- Stage (3) – Shape identification
 - Rectangle:
 - 4 vertices
 - All angles are ~ 90 degrees apart
 - Checked using θ
 - Triangle
 - 3 vertices



FYI: Digit Recognition

- Contour is used in the digit recognition program

1AB

Detection of other shapes

- Hough Transform can be generalized to find other shapes like circles and ellipses.
- However, the computational complexity increases
 - More computational time is required

Hough circle transform

- Equation of circle:

$$(x - a)^2 + (y - b)^2 = R^2$$

- (a, b) is the center of the circle
- R is the radius of the circle

- Parameters: a, b, R


- Rewriting the equation:

$$(a - x)^2 + (b - y)^2 = R^2$$

Circle with known R

$$(a - x)^2 + (b - y)^2 = R^2$$


- Parameters: a, b
- Points (edge) in picture \rightarrow circle in parameter space
- Points in parameter space \rightarrow circle in picture
- Like line detection, discretization and voting are used
- Example:



Circle with unknown R

- Parameters: a, b, R
- Points (edge) in picture \leftrightarrow conical surface in parameter space
- Additional computational time

Pixelization
Voxelization




Application: Extraction of Blastomere

- Embryo Biopsy
- 8-cell stage
- Tracking of the extraction of blastomere for biopsy



First Frame



Circle detection

Blastomere tracking